



Microservices & docker: from theory to practice

- Tom Van Cutsem, PhD and Nico Janssens, PhD
- Bell Labs, Application platforms and software systems lab
- December 1st, 2016


Context




What drives microservices?



You're in good company



What are microservices?



“SOA done right”




-- *James Lewis and Martin Fowler*

Microservices: characteristics



Componentization
via services




Smart endpoints,
dumb pipes

Microservices: characteristics



Products,
not projects

“you build it, you run it”




Organized around
business capabilities

Microservices: characteristics




Products,
not projects

“you build it, you run it”



Organized around
business capabilities

Microservices: organize around business services



"Any organization that designs a system ... will inevitably produce a design whose structure is a copy of the organization's communication structure."
-- Melvin Conway, 1968

(Source: Martin Fowler)

Microservices: characteristics




Decentralized
Governance



Decentralized
Data Management


Large codebases seem to auto-evolve into microservices

War stories from large web companies



(Source: highscalability.com, 2015)

Microservices: characteristics




Infrastructure
Automation

“IT is an API”





Design for Failure



Evolutionary
Design

Microservices & DevOps culture

- Need to be able to provision infrastructure *fast*
- Containerize services (Docker)
- Container orchestration (Swarm, Kubernetes, Rancher, Mesos, ...)
- Teams maintain their own services in production



Microservices: risks

Independent services



Service boundaries
not easy to change

End-to-end testing/debugging
more difficult

Distributed systems
challenges

Design for Failure



Investment in monitoring
tools

Operational complexity

Technological Diversity




Strong and diverse
skill set


Case study: ShowMe

ShowMe: location-based video sharing

Discover or share what's up near a location of interest






LIVE PERSONAL
CONTENT




LOCATION AS
SHARED CONTEXT

ShowMe: location-based video sharing


Prototype app + experience movie




Functional architecture



Technology stack



Service interfaces



Lessons learned

- Multiple teams working on independent subsystems = highly productive
 - Different goals
 - Different skillsets
 - Different release schedules
 - Less conflicts
- Testing and debugging of the overall system was a pain
- We didn't sufficiently invest in tooling and automation
 - Manual configuration and set-up
 - Infrastructure not set up to host multiple versions of the app
 - No cross-service unit testing infrastructure


Lessons learned

- Micro-service architecture = distributed system
 - Deal with asynchrony, failure, latency, keeping data consistent across databases
 - Interfaces between services are implicit, not checked by compiler.
- Testing services in isolation is not enough
 - Focus is on monitoring and detecting anomalies more than on thorough testing before deployment
- Deployment is much more complicated
 - Fine-grained orchestration and configuration
 - Each service needs clustering, monitoring, load-balancing, ...
 - Variety of runtimes and databases requires larger skill set to tweak, deploy, maintain
 - To do microservices right, should keep old and new versions of the service running side-by-side

Case study: instadash

Instdash app

Real-time fleet tracking



GPS receiver



OBD via CAN bus



Dashcam




On-board Unit




2 real cars,
10 hours footage
400 virtual cars




Instadash: functional architecture




Instdash: functional architecture




Instdash: technology stack



Instdash: service interfaces



Microservice communication patterns




REST/HTTP is well-understood

HTTP support is ubiquitous

JSON as data model is a natural fit

Text-based protocol overheads

Microservice communication patterns



Fast, often binary encoding

Built-in schema support

Firewall issues, less ubiquitous

Need an additional discovery service

Public



JSON-RPC



Microservice communication patterns




Decoupling between components
(bus handles both discovery and routing)

More complex, beware
bottleneck


Public




Communication patterns



Communication patterns: point-to-point




Communication patterns: publish-subscribe




- Use AMQP's rich routing semantics via topic exchanges
- Messages represent events (JSON payload)

Communication patterns: work queueing




Monitoring: our approach

- Used Riemann as central dashboard and event monitoring server
- Client libraries for a variety of programming languages (remember: polyglot)
- Each microservice regularly reports service-specific statistics
- Each host machine also reports generic resource statistics



Monitoring



Monitoring: dashboards




Lessons learned

- Message bus as **central broker** had many advantages
 - Solved service discovery (all components need to know the broker, not each other)
 - Queueing makes services more robust to failover
 - Message bus dashboard gave a wealth of system information about communication patterns, message rates, etc.
 - But: can quickly become a bottleneck: proper configuration and tuning was key
 - Also: all components needed hardening to e.g. auto-reconnect when broker went down
- Use external **configuration files** that can be generated or templated from a central place
- Use **schema validation** to catch bugs faster (e.g. JSON-Schema, Protobufs, AVRO, ...)
- **Monitoring** was essential to see what's going on
- **Dockerizing** services was key to getting this system going (20+ processes)

PART III

On MicroServices, Docker and
DevOps


Context




Docker Containers

Efficiency

Lightweight application isolation → very low performance overhead



Virtual Machines



Containers

source: <https://blog.jayway.com/2015/03/21/a-not-very-short-introduction-to-docker/>

Docker Containers

Programmability

Container programming → Dockerfile

```
FROM ubuntu:16.04
MAINTAINER Sven Dowideit <SvenDowideit@docker.com>

RUN apt-get update && apt-get install -y openssh-server
RUN mkdir /var/run/sshd
RUN echo 'root:screencast' | chpasswd
RUN sed -i 's/PermitRootLogin prohibit-password/PermitRootLogin yes/' /etc/ssh/sshd_config

# SSH login fix. Otherwise user is kicked off after login
RUN sed 's@session\s*required\s*pam_loginuid.so@session optional pam_loginuid.so@g' -i /etc/pam.d/sshd

ENV NOTVISIBLE "in users profile"
RUN echo "export VISIBLE=now" >> /etc/profile

EXPOSE 22
CMD ["/usr/sbin/sshd", "-D"]
```

Docker Containers

Active eco-system



Docker Containers

Portability


The matrix from hell

		?	?	?	?	?	?	?
			?	?	?	?	?	?
		?	?	?	?	?	?	?
		?	?	?	?	?	?	?
		?	?	?	?	?	?	?
		?	?	?	?	?	?	?

Docker Containers

Flexibility


Build, ship and run any app, anywhere [docker]



Docker Containers

Demo

DevOps @10k feet



DevOps

CALMS

Culture

- Promotes collaborative and open culture between Dev and Ops
- Embrace change and experimentation

Automation

- Automate whenever possible
- CI/CD, Infrastructure as Code, ...

Lean

- Focus on producing value for the end-user
- Small size batches, higher release cycles

Measurement

- Measure everything all the time and use this info to improve/refine cycles
- Show the improvement

Sharing

- Open information sharing – experiences, successes, failures, etc.
- Collaboration & communication – learn from each other (Kanban board, IM, wiki)

Moving away from traditional telco service design

Operational costs pressures push Telcos to virtualize environments while preserving **non-functional requirements**



- 5 nines availability
- Reliability
- Performance and response times

Moving away from traditional telco service design

Additional **non-functional requirements** to take into account

- Scalability
- Elasticity
- Agility
- Operability and portability



Low overhead


Portability

Micro-service architectures


Active eco-system + public image registries

Facilitates DevOps methodology

Bell Labs Projects: New Home/IoT Service Platform



Bell Labs Projects: Bandwidth Optimized Streaming Analytics




B. Theeten and N. Janssens, CHive:
Bandwidth Optimized Continuous
Querying in Distributed Clouds. IEEE
Transactions on Cloud Computing 3
(2), pp. 219-232, 2015.


Bell Labs Projects: New Communication Service

Key Goal: Simplify interactions among people, machines, and their environments

- From transaction-oriented Web model to persistent conversations
- Uniform interaction model for people, machines, and objects
- Rich context-based communications and collaboration




Micro-service chat architecture



Initial production design

20 node cluster with RANCHER and DOCKER



Evaluation

MicroServices

Rapid and independent evolution (lifecycle management) ✓

Use the right tool for the job ✓

Decentralized governance and data management ✓

Evaluation

Docker


Low overhead ✓

Portability ✓


Micro-service architectures ✓

Active eco-system with public image registries ✓

Facilitates DevOps methodology ✓



toolbox



compose



registry



Evaluation

Docker

Docker lifecycle management

- Don't forget to clean old containers and dangling images
- For non-trivial lifecycle mgmt and production environments, rely on other tools
 - compose, swarm, kubernetes, mesos+marathon/chronos, saltstack, terraform, etc.

Dockerfiles

- Think carefully how to structure your Dockerfiles (across Dockerfiles)
 - Each line in a Docker file is a separate image layer, which by default will be cached (exceptions!)
- Order from generic/stable commands to specific/unstable commands
 - Use explicit version tagging for all installed packages (consistency across future builds)
 - Avoid unnecessary layers & packages → smartly combine commands

Performance when sharing host resources (e.g. when using bridge network)

No need to dockerize all your services ...

Evaluation

Docker

Application packaging → KISS!

- Containers are not VMs, but application environments
- Don't try to stuff too many background services inside each container (sshd, logging, etc.)
- Don't install build tools (e.g. gcc) without good reason → use build containers for that!

Data storage

- Try to avoid storing (all) data inside the application containers
 - Containers should be as much as possible easily replaceable
- Use key-value stores (etcd), DBs (mysql), data containers or host-volumes (-v)

Security

Networking

Background reading and references

- Martin Fowler's article (must read): <http://martinfowler.com/articles/microservices.html>
- Community site: <http://microservices.io/>
- A. Cockcroft (prev. Netflix lead engineer) on migrating to micro-services:
<http://www.infoq.com/presentations/migration-cloud-native>
- Insightful blogs:
 - <http://www.tigerteam.dk/2014/micro-services-its-not-only-the-size-that-matters-its-also-how-you-use-them-part-1/>
 - <http://gomorpheus.com/blog/2014-10-24-the-new-reality-microservices-apply-the-internet-model-to-app-development>
 - A critical note: <http://contino.co.uk/microservices-not-a-free-lunch/>
 - <http://highscalability.com/blog/2015/12/1/deep-lessons-from-google-and-ebay-on-building-ecosystems-of.html>
- Colossus (Tumblr Engineering Blog): <http://engineering.tumblr.com/post/102906359034/colossus-a-new-service-framework-from-tumblr>
- Finagle (Twitter Engineering Blog): <https://blog.twitter.com/2011/finagle-a-protocol-agnostic-rpc-system>

NOKIA

Copyright and confidentiality

The contents of this document are proprietary and confidential property of Nokia. This document is provided subject to confidentiality obligations of the applicable agreement(s).

This document is intended for use of Nokia's customers and collaborators only for the purpose for which this document is submitted by Nokia. No part of this document may be reproduced or made available to the public or to any third party in any form or means without the prior written permission of Nokia. This document is to be used by properly trained professional personnel. Any use of the contents in this document is limited strictly to the use(s) specifically created in the applicable agreement(s) under which the document is submitted. The user of this document may voluntarily provide suggestions, comments or other feedback to Nokia in respect of the contents of this document ("Feedback").

Such Feedback may be used in Nokia products and related specifications or other documentation. Accordingly, if the user of this document gives Nokia Feedback on the contents of this document, Nokia may freely use, disclose, reproduce, license, distribute and otherwise commercialize the feedback in any Nokia product, technology, service, specification or other documentation.

Nokia operates a policy of ongoing development. Nokia reserves the right to make changes and improvements to any of the products and/or services described in this document or withdraw this document at any time without prior notice.

The contents of this document are provided "as is". Except as required by applicable law, no warranties of any kind, either express or implied, including, but not limited to, the implied warranties of merchantability and fitness for a particular purpose,

are made in relation to the accuracy, reliability or contents of this document. **NOKIA SHALL NOT BE RESPONSIBLE IN ANY EVENT FOR ERRORS IN THIS DOCUMENT or for any loss of data or income or any special, incidental, consequential, indirect or direct damages howsoever caused, that might arise from the use of this document or any contents of this document.**

This document and the product(s) it describes are protected by copyright according to the applicable laws.

Nokia is a registered trademark of Nokia Corporation. Other product and company names mentioned herein may be trademarks or trade names of their respective owners.