

Designing “least-authority” JavaScript apps

Tom Van Cutsem
DistriNet KU Leuven



tvcutsem.github.io



be.linkedin.com/in/tomvc



github.com/tvcutsem



x.com/tvcutsem



@tvcutsem@techhub.social

Web application security

same-origin policy

certificate pinning

OAuth

cookies

content security policy

CSRF

HSTS

html sanitization



OWASP

Open Web Application
Security Project

A **software engineering view** of Web application security

~~same-origin policy~~

modules

~~certificate pinning~~

functions

~~OAuth~~

~~cookies~~

encapsulation

dependencies

~~content security policy~~

immutability

~~CSRF~~

dataflow

~~HSTS~~

~~html sanitization~~

isolation

A software engineering view of Web application security

“Security is just the extreme of Modularity”

- Mark S. Miller
(Chief Scientist, Agoric)




Modularity: avoid needless dependencies (to prevent bugs)

Security: avoid needless vulnerabilities (to prevent exploits)

The CIA triad from an application security perspective

- **Confidentiality** (a.k.a. Secrecy): No one can infer information they are not supposed to know. Confidentiality usually rests on **cryptography** to keep information secret.
 - Example violation: “*Bob learns how much money Alice has in her bank account*”
 - Example threat: side channel attack.
- **Integrity** (a.k.a. Safety): No “bad” things happen. Integrity usually rests on **access control** determining what agents can cause what effects.
 - Example violation: “*Bob steals Alice’s money*”
 - Example threat: confused deputy attack.
- **Availability** (a.k.a. Liveness): “Good” things continue to happen.
 - Example violation: “*Bob prevents Alice from spending her money as she wants*”
 - Example threat: a denial of service attack.




(Image source: Nikander, Jussi & Manninen, Onni & Laajalahti, Mikko. (2020). Requirements for cybersecurity in agricultural communication networks. Computers and Electronics in Agriculture.)

(Source: Miller, “A Taxonomy of Security Issues”, 2021, <https://agoric.com/blog/technology/a-taxonomy-of-security-issues>)

The CIA triad from an application security perspective

- **Confidentiality** (a.k.a. Secrecy): No one can infer information they are not supposed to know. Confidentiality usually rests on **cryptography** to keep information secret.
 - Example violation: “*Bob learns how much money Alice has in her bank account*”
 - Example threat: side channel attack.
- **Integrity** (a.k.a. Safety): No “bad” things happen. Integrity usually rests on **access control** determining what agents can cause what effects.
 - Example violation: “*Bob steals Alice’s money*”
 - Example threat: confused deputy attack.
- **Availability** (a.k.a. Liveness): “Good” things continue to happen.
 - Example violation: “*Bob prevents Alice from spending her money as she wants*”
 - Example threat: a denial of service attack.






Our focus




(Image source: Nikander, Jussi & Manninen, Onni & Laajalahti, Mikko. (2020). Requirements for cybersecurity in agricultural communication networks. Computers and Electronics in Agriculture.)

(Source: Miller, “A Taxonomy of Security Issues”, 2021, <https://agoric.com/blog/technology/a-taxonomy-of-security-issues>)

Application integrity & access control



Application integrity (safety): going beyond OS process isolation





This Lecture

- Part I: **why module isolation** is critical to modern JavaScript applications
- Part II: the **Principle of Least Authority**, by example (in JavaScript)
- Part III: safely composing modules using **least-authority patterns**

Part I

Why module isolation is critical to modern JavaScript applications

JavaScript is no longer just about the Web. Used widely across *all* tiers.



GraalVM™



Embedded


Mobile

Desktop/Native

Server

Database

Modern JavaScript applications are built from thousands of modules



2,000,000 modules on NPM


“The average modern web application has over 1000 modules [...] **97% of the code in a modern web application comes from npm**. An individual developer is responsible only for the final 3% that makes their application unique and useful.”

(source: *npm blog*, December 2018)

(source: modulecounts.com, Nov 2022)



Composing modules: it's all about **trust**

It is exceedingly common to run code you don't know or trust in a common environment




What can happen when a module goes **rogue**?

It is exceedingly common to run code you don't know or trust in a common environment



What can happen when a module goes **rogue**?



```
<script src="http://evil.com/ad.js">
```


The New York Times @nytimes

Attn: NYTimes.com readers: Do not click pop-up box warning about a virus -- it's an unauthorized ad we are working to eliminate.

17 7:54 PM - Sep 13, 2009

[See The New York Times's other Tweets](#) >

What can happen when a module goes **rogue**?



Check your repos... Crypto-coin-stealing code sneaks into fairly popular NPM lib (2m downloads per week)

Node.js package tried to plunder Bitcoin wallets

By Thomas Claburn in San Francisco 26 Nov 2018 at 20:58 49 □ SHARE ▾

A screenshot of a GitHub commit showing a block of malicious JavaScript code. The code appears to be part of a plugin for a carousell component, containing logic to steal Bitcoin wallets. The code includes variables like \$target, options, slideIndex, and target, along with various jQuery-like methods and conditions.

```
$(this).attr('data-target')) ||  
  this.$target = $($this.attr('data-target') || '') // st  
  href.replace(.*(?=#[^\s]+$/), '')  
  if ($target.hasClass('carousell')) return  
  options = $.extend({}, $target.data(), $  
  slideIndex = $this.attr('data-slide-to')  
  (slideIndex) options.interval = false  
  plugin.call(target, options)  
  if (slideIndex) {  
    target.data('bs.carou...
```

(source: [theregister.co.uk](https://www.theregister.co.uk/2018/11/26/crypto-coin-stealing_code_sneaks_into_fairly_popular_NPM_lib/))

These are examples of software supply chain attacks

Software Supply Chain Security | August 18, 2022

6 reasons app sec teams should shift gears and go beyond legacy vulnerabilities

 BLOG AUTHOR
John P. Mello Jr., Freelance technology writer. [READ MORE...](#)

[!\[\]\(f9640f02d06b604c50c1e3a558a31388_img.jpg\)](#) [!\[\]\(d9d98187a25260cd1316acead95958b1_img.jpg\)](#) [!\[\]\(5e91ce58ee8e9ab2222ea2dcf7128599_img.jpg\)](#) [!\[\]\(0fee90055c33b690097ce6e49c052bde_img.jpg\)](#)



With software supply chain attacks surging, dev and application security teams should shift gears from legacy vulnerabilities to open-source repos, DevOps tools, and software tampering.

1. Trusting code within the supply chain has become problematic

Many tools designed to help secure software-development pipelines focus on rating the projects, programmers, and open-source components and their maintainers. However, recent events—such as the emergence of the “protestware” that changed the node.ipc open source software for political reasons or the hijacking of the popular ua-parser-js project by cryptominer—underscore that seemingly secure projects can be compromised, or otherwise pose security risks to organizations. ”

Tomislav Peričin, co-founder and chief software architect at ReversingLabs, noted how [in the case of SolarWinds](#), the trusted source was pushing infected software. Catching those kinds of mistakes requires a focus on how code behaves, regardless of where it came from.

“As long as we keep ignoring the core of the problem – which is how do you trust code – we are not handling software supply chain security.”

—[Tomislav Peričin](#)

(Source: <https://develop.secure.software/6-reasons-software-security-teams-need-to-go-beyond-vulnerability-response>, august 2022)

Increasing awareness

Great tools, but address the symptoms, not the root cause

npm security advisories

Security advisories		
Advisory	Date of advisory	Status
Cross-Site Scripting bootstrap-select severity high	May 20th, 2020	status patched
Cross-Site Scripting @toast-ui/editor severity high	May 20th, 2020	status patched
Cross-Site Scripting jquery severity moderate	Apr 30th, 2020	status patched

npm audit

```
==== npm audit security report ====
# Run `npm install chokidar@2.8.3` to resolve 1 vulnerability
SEMVER WARNING: Recommended action is a potentially breaking change
Low          Prototype Pollution
Package      deep-extend
Dependency of chokidar
Path         chokidar > fsevents > node-pre-gyp > rc > deep-extend
More info    https://nodesecurity.io/advisories/612
```

GitHub security alerts

28 commits 1 branch 0 packages 2 releases 2 contributors MIT

We found potential security vulnerabilities in your dependencies.

Only the owner of this repository can see this message.

View security alerts

Snyk vulnerability DB

snyk Test Features Vulnerability DB Blog Partners Pricing Docs About Log In Sign Up

Vulnerability DB > npm > lodash

Prototype Pollution
Affecting lodash package, ALL versions

Report new vulnerabilities

Do your applications use this vulnerable package? Test your applications

Overview
lodash is a modern JavaScript utility library delivering modularity, performance, & extras.
Affected versions of this package are vulnerable to Prototype Pollution. The function `zipObjectDeep` can be tricked into adding or modifying properties of the Object prototype. These properties will be present on all objects.


CVSS SCORE 6.3 MEDIUM SEVERITY

ATTACK VECTOR Network ATTACK COMPLEXITY Low

PRIVILEGES REQUIRED Low USER INTERACTION None

Avoiding interference is the name of the game

- Shield important resources/APIs from modules that don't need access
- Apply **Principle of Least Authority** (POLA) to application design




Part II

The Principle of Least Authority, by example (in JavaScript)


Principle of Least Authority (POLA)

- A module should only be given the authority it needs to do its job, and nothing more




What is “authority” in a JavaScript app?

- Authority is linked to resources represented as objects (or functions)
- Objects can hold references (“pointers”) to resource objects
- The authority to use a resource is expressed by calling a method/function on a reference



Delegating authority == sharing references, under the right assumptions



(source: Miller et al. "Capability myths demolished", 2003)

Example: Alice wants to give Bob access to Carol, and *only* to Carol:

```
// alice calls:  
bob.foo(carol)
```

Assumptions:

- Pointers (references) are unforgeable
✓ JavaScript is **memory-safe**
- Pointers (references) can be privately stored
✓ JavaScript supports **hiding** access to private state through scoping rules
- There is no global mutable state
⚠ Ensure that all exported objects/functions in a module are **immutable**
- There is no undeniable ("ambient") authority
⚠ Need to load each module in its **own**, initially empty, **global environment**


Running example: apply POLA to a basic shared log

We would like Alice to only **write** to the log, and Bob to only **read** from the log.

```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}

let log = new Log();
alice(log);
bob(log);
```




Running example: apply POLA to a basic shared log

If Bob goes rogue, what could go wrong?

```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}

let log = new Log();
alice(log);
bob(log);
```



Bob has way too much authority!

If Bob goes rogue, what could go wrong?

```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}


let log = new Log();
alice(log);
bob(log);
```

```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")

// Bob can delete the entire log
log.read().length = 0


// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can replace the Array built-ins
Array.prototype.push = function(msg) {
  console.log("I'm not logging anything");
}
```



How to solve “prototype poisoning” attacks?

Load each module in its own environment,
with its own set of “primordial” objects



```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}

let log = new Log();
alice(log);
bob(log);

// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")


// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can replace the Array built-ins
Array.prototype.push = function(msg) {
  console.log("I'm not logging anything");
}
```


Prerequisite: isolating JavaScript modules

- Today: JavaScript offers no standard way to isolate a module (load it in a separate environment)
- Lots of host-specific isolation mechanisms, but non-portable and ill-defined:
 - **Web Workers**: no shared memory, can only communicate using message-passing
 - **iframes**: mutable primordials, “identity discontinuity”
 - **nodejs vm module**: same issues



ShadowRealms (ECMA TC39 Stage 2 proposal)


Intuitions: “iframe without DOM”, “principled version of node’s `vm` module”



* Primordials: built-in objects like Object, Object.prototype, Array, Function, Math, JSON, etc.

Compartments (ECMA TC39 Stage 1 proposal)

Each Compartment has its own global object but shared (immutable) primordials.



* Primordials: built-in objects like Object, Object.prototype, Array, Function, Math, JSON, etc.

Hardened JavaScript is a secure subset of standard JavaScript



Key idea: code running in hardened JS can only affect the outside world through objects (capabilities) explicitly granted to it from outside.

(inspired by the diagram at <https://github.com/Agoric/Jessie>)

Hardened JavaScript: some history

- Google develops a project called “**Caja**” for **safe embedding** of dynamic web content (JavaScript scripts) in web pages
- Attempts are made to **standardize** core features that enable secure sandboxing as “**Secure ECMAScript**” (SES) at ECMA TC39
- Standardisation process got stalled, but work continued on a modified node.js runtime called “**endo**”, supporting SES on the server
- A company called Agoric **rebrands** SES to “**Hardened JavaScript**”, works with Moddable and Metamask on implementation and tooling
- HardenedJS is **used by several companies** to isolate JavaScript modules for IoT (Moddable), Web3 (Agoric), SaaS (Salesforce), ...



LavaMoat

- CLI tool that puts each package dependency into its own hardened JS sandbox environment
- Auto-generates config file indicating authority needed by each package
- Plugs into build tools like Webpack and Browserify



```
npm install -D lavamoat  
npx lavamoat app.js --autopolocy
```



<https://github.com/LavaMoat/lavamoat>




```
"stream-http": {  
    "globals": {  
        "Blob": true,  
        "MSStreamReader": true,  
        "ReadableStream": true,  
        "VBArrary": true,  
        "XDomainRequest": true,  
        "XMLHttpRequest": true,  
        "fetch": true,  
        "location.protocol.search": true  
    },  
    "packages": {  
        "buffer": true,  
        "builtin-status-codes": true,  
        "inherits": true,  
        "process": true,  
        "readable-stream": true,  
        "to-arraybuffer": true,  
        "url": true,  
        "xtend": true  
    }  
},
```

LavaMoat enables more focused security reviews

Exposure to package dependencies
without LavaMoat sandboxing



Exposure to package dependencies
with LavaMoat sandboxing



<https://github.com/LavaMoat/lavamoat>

Bonus: avoiding unwanted post-install scripts

- Package managers like `npm` allow packages to run install scripts
- A compromised dependency can exploit this to run code as part of your project installation script
- Lavamoat's `allow-scripts` tool configures your project to disable running install scripts by default
- Edit allowed packages in `package.json`
- New install scripts entering your dependency tree will no longer run automatically unless approved




```
npm install -D @lavamoat/allow-scripts
npx --no-install allow-scripts auto

// in package.json
{
  "lavamoat": {
    "allowScripts": {
      "keccak": true,
      "core-js": false
    }
  }
}
```

<https://www.npmjs.com/package/@lavamoat/allow-scripts>

Back to our example

With Alice and Bob's code running in their own Compartment, we mitigate the poisoning attack



```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}

let log = new Log();
alice(log);
bob(log);

// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")


// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can replace the Array built-ins
Array.prototype.push = function(msg) {
  console.log("I'm not logging anything");
}
```

One down, three to go

POLA: we would like Alice to only write to the log, and Bob to only read from the log.



```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}

let log = new Log();
alice(log);
bob(log);
```

```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")

// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}
```

Make the log's interface **tamper-proof**

Object.freeze makes property bindings (not their values) immutable

```
import * as alice from "alice.js";
import * as bob from "bob.js";

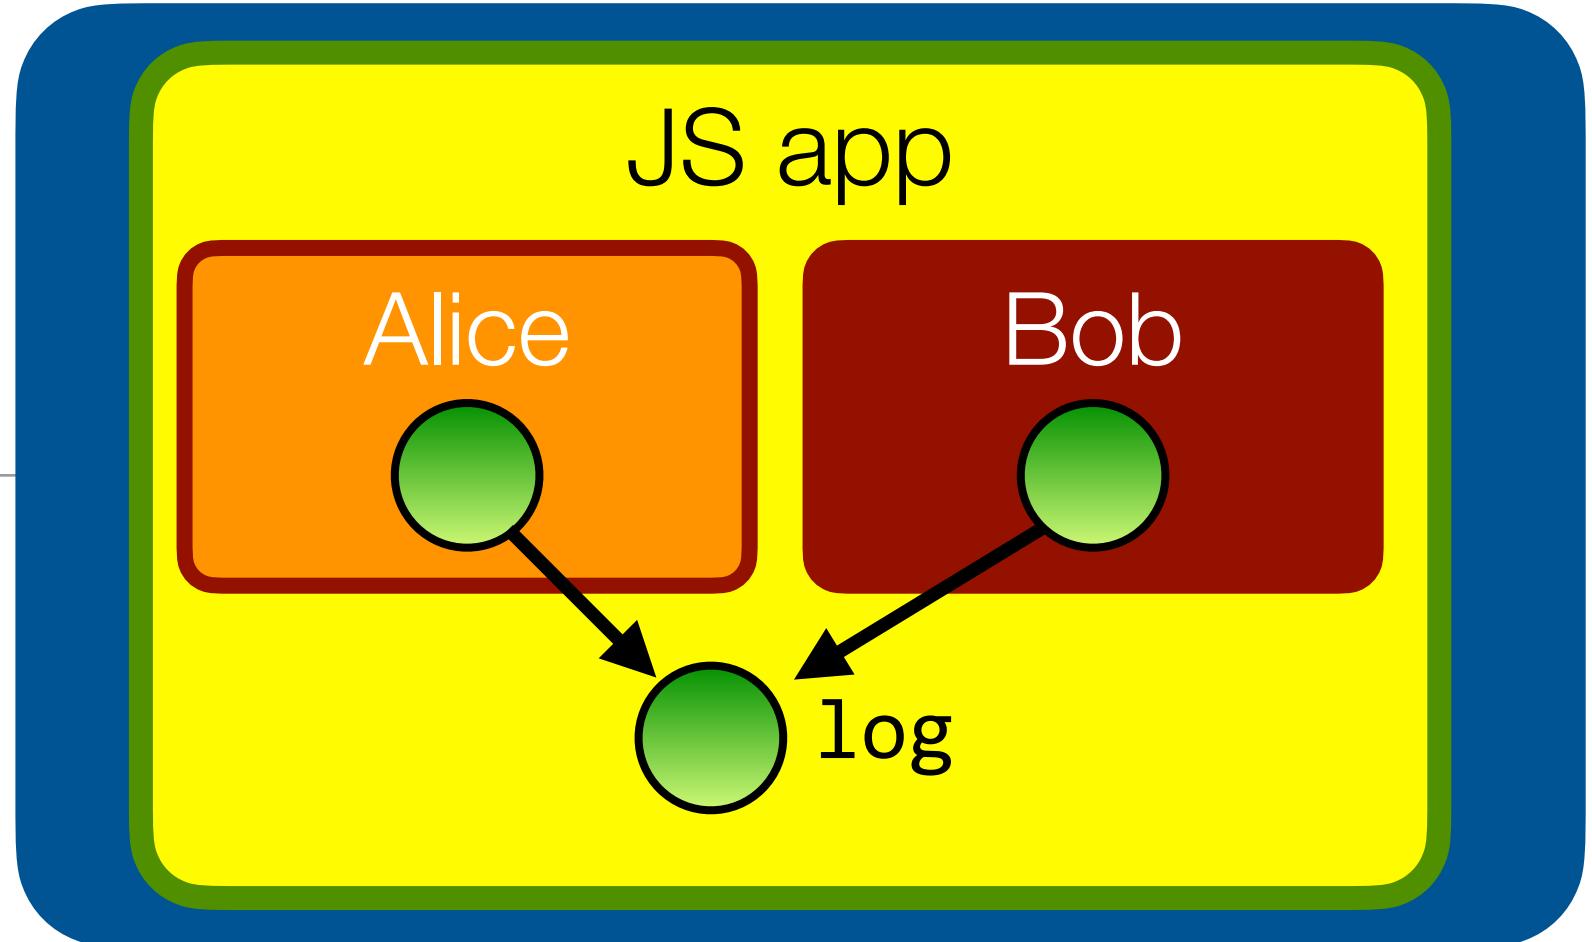
class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}

let log = Object.freeze(new Log());
alice(log);
bob(log);
```

```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")


// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}
```



Make the log's interface tamper-proof. Oops.

Functions are mutable too. Freeze doesn't recursively freeze the object's functions.



```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}

let log = Object.freeze(new Log());
alice(log);
bob(log);
```

```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")

// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```

Make the log's interface tamper-proof

Hardened JavaScript provides a `harden` function that “deep-freezes” an object

```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}


let log = harden(new Log());
alice(log);
bob(log);
```

```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")


// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```



Two down, two to go



```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}

let log = harden(new Log());
alice(log);
bob(log);
```


```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")

// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```

Two down, two to go



```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return this.messages_; }
}

let log = harden(new Log());
alice(log);
bob(log);
```

```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")

// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```

Don't share access to mutable internals

- Modify `read()` to return a copy of the mutable state.
- Even better would be to use a more efficient copy-on-write or “immutable” data structure (see immutable-js.com)

```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return [...this.messages_]; }
}


let log = harden(new Log());
alice(log);
bob(log);
```

```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")


// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```



Three down, one to go



```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return [...this.messages_]; }
}

let log = harden(new Log());
alice(log);
bob(log);
```

```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")

// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```


Three down, one to go

- Recall: we would like Alice to only write to the log, and Bob to only read from the log.
- Bob receives too much authority. How to limit?

```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return [...this.messages_]; }
}

let log = harden(new Log());
alice(log);
bob(log);
```



```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")

// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```

Pass only the authority that Bob needs.

Just pass the write function to Alice and the read function to Bob.

```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return [...this.messages_]; }
}


let log = new Log();
let read = harden(() => log.read());
let write = harden((msg) => log.write(msg));
alice(write);
bob(read);

// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")


// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```



Success! We thwarted all of Evil Bob's attacks.



```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return [...this.messages_]; }
}

let log = new Log();
let read = harden(() => log.read());
let write = harden((msg) => log.write(msg));
alice(write);
bob(read);

// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")

// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```

Is there a better way to write this code?

The burden of correct use is on the *client* of the class. Can we avoid this?

```
import * as alice from "alice.js";
import * as bob from "bob.js";

class Log {
  constructor() {
    this.messages_ = [];
  }
  write(msg) { this.messages_.push(msg); }
  read() { return [...this.messages_]; }
}


let log = new Log();
let read = harden(() => log.read());
let write = harden((msg) => log.write(msg));
alice(write);
bob(read);
```

```
// in bob.js
// Bob can just write to the log
log.write("I'm polluting the log")

// Bob can delete the entire log
log.read().length = 0

// Bob can replace the 'write' function
log.write = function(msg) {
  console.log("I'm not logging anything");
}

// Bob can still modify the write function
log.write.apply = function() { "gotcha" };
```




Use the **Function as Object** pattern

- A record of closures hiding state is a fine representation of an object of methods hiding instance vars
- Pattern long advocated by Doug Crockford instead of using classes or prototypes


```
class Log {  
  constructor() {  
    this.messages_ = [];  
  }  
  write(msg) { this.messages_.push(msg); }  
  read() { return [...this.messages_]; }  
}
```

```
let log = new Log();  
let read = harden(() => log.read());  
let write = harden((msg) => log.write(msg));  
alice(write);  
bob(read);
```




```
function makeLog() {  
  const messages = [];  
  function write(msg) { messages.push(msg); }  
  function read() { return [...messages]; }  
  return harden({read, write});  
}
```

```
let log = makeLog();  
alice(log.write);  
bob(log.read);
```



(See also <https://martinfowler.com/bliki/FunctionAsObject.html>)

Use the Function as Object pattern




```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
  const messages = [];
  function write(msg) { messages.push(msg); }
  function read() { return [...messages]; }
  return harden({read, write});
}

let log = makeLog();
alice(log.write);
bob(log.read);
```

What if Alice and Bob need more authority?


If over time we want to expose more functionality to Alice and Bob, we need to refactor all of our code.



```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
  const messages = [];
  function write(msg) { messages.push(msg); }
  function read() { return [...messages]; }
  return harden({read, write});
}

let log = makeLog();
alice(log.write);
bob(log.read);
```




```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
  const messages = [];
  function write(msg) { messages.push(msg); }
  function read() { return [...messages]; }
  function size() { return messages.length(); }
  return harden({read, write, size});
}

let log = makeLog();
alice(log.write, log.size);
bob(log.read, log.size);
```

Expose distinct authorities through facets


Easily deconstruct the API of a single powerful object into separate interfaces by nesting objects



```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
  const messages = [];
  function write(msg) { messages.push(msg); }
  function read() { return [...messages]; }
  function size() { return messages.length(); }
  return harden({read, write, size});
}

let log = makeLog();
alice(log.write, log.size);
bob(log.read, log.size);
```



```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
  const messages = [];
  function write(msg) { messages.push(msg); }
  function read() { return [...messages]; }
  function size() { return messages.length(); }
  return harden({
    reader: {read, size},
    writer: {write, size}
  });
}


let log = makeLog();
alice(log.writer);
bob(log.reader);
```

Demo

<https://github.com/tvcutsem/lavamoat-demo>

End of Part II: recap


- Modern JS apps are composed from many modules. You can't trust them all.
- Traditional security boundaries don't exist between modules. Compartments add basic isolation.
- **Isolated modules must still interact!**
- **Compose** functionality from untrusted modules in a **least-authority** manner
- This can be done via **reusable programming patterns** that rely on object-capability security



Part III


Safely composing modules using least-authority patterns

Design Patterns (“Gang of Four”, 1994)



- Visitor
- Factory
- Observer
- Singleton
- State
- ...

Design Patterns for **robust composition** (Mark S. Miller, 2006)




- Facets
- Taming
- Caretaker
- Membrane
- Sealer/unsealer pair
- ...

<http://www.erights.org/talks/thesis/markm-thesis.pdf>

Recall: the Principle of Least Authority (POLA)

- A module should only be given the authority it needs to do its job, and nothing more



Further limiting Bob's authority


We would like to give Bob only **temporary** read access to the log.

```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
    const messages = [];
    function write(msg) { messages.push(msg); }
    function read() { return [...messages]; }
    return harden({read, write});
}

let log = makeLog();

alice(log.write);
bob(log.read);
```




Use **caretaker** to insert access control logic

We would like to give Bob only **temporary** read access to the log.

```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
    const messages = [];
    function write(msg) { messages.push(msg); }
    function read() { return [...messages]; }
    return harden({read, write});
}

let log = makeLog();
let [rlog, revoke] = makeRevokableLog(log);
alice(log.write);
bob(rlog.read);
```



Use **caretaker** to insert access control logic


We would like to give Bob only **temporary** read access to the log.

```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
    const messages = [];
    function write(msg) { messages.push(msg); }
    function read() { return [...messages]; }
    return harden({read, write});
}

let log = makeLog();
let [rlog, revoke] = makeRevokableLog(log);
alice(log.write);
bob(rlog.read);

// to revoke Bob's access:
revoke();
```




Use **caretaker** to insert access control logic

```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
  const messages = [];
  function write(msg) { messages.push(msg); }
  function read() { return [...messages]; }
  return harden({read, write});
}

let log = makeLog();
let [rlog, revoke] = makeRevokableLog(log);
alice(log.write);
bob(rlog.read);

// to revoke Bob's access:
revoke();
```



```
function makeRevokableLog(log) {
  function revoke() { log = null; }
  let proxy = {
    write(msg) { log.write(msg); }
    read() { return log.read(); }
  };
  return harden([proxy, revoke]);
}
```


A caretaker is just a proxy object

```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
  const messages = [];
  function write(msg) { messages.push(msg); }
  function read() { return [...messages]; }
  return harden({read, write});
}


let log = makeLog();
let [rlog, revoke] = makeRevokableLog(log);
alice(log.write);
bob(rlog.read);

// to revoke Bob's access:
revoke();
```



```
function makeRevokableLog(log) {
  function revoke() { log = null; }
  let proxy = {
    write(msg) { log.write(msg); }
    read() { return log.read(); }
  };
  return harden([proxy, revoke]);
}
```

A caretaker is just a proxy object



```
import * as alice from "alice.js";
import * as bob from "bob.js";

function makeLog() {
  const messages = [];
  function write(msg) { messages.push(msg); }
  function read() { return [...messages]; }
  return harden({read, write});
}


let log = makeLog();
let [rlog, revoke] = makeRevokableLog(log);
alice(log.write);
bob(rlog.read);

// to revoke Bob's access:
revoke();
```

```
function makeRevokableLog(log) {
  function revoke() { log = null; }
  let proxy = {
    write(msg) { log.write(msg); }
    read() { return log.read(); }
  };
  return harden([proxy, revoke]);
}
```


Taming is the process of restricting access to powerful APIs

- Expose powerful objects through restrictive proxies to third-party code
- E.g. Alice might give Bob read-only access to a specific subdirectory of her file system



Taming is the process of restricting access to powerful APIs


Potential hazard: the taming proxy must ensure it does not “leak” privileged access to host resources through the tamed API (e.g. through return values)



Taming is the process of restricting access to powerful APIs


The **solution** is to transitively apply the proxy pattern to return values as well. This pattern is called a “**membrane**”

Deep dive blog post at tvcutsem.github.io/membranes




Least-authority patterns are used in industry

Example: how Google Caja uses taming to restrict access to the browser DOM



Google Caja



(source: Google Caja documentation: <https://developers.google.com/caja/docs/about>)

Least-authority patterns are used in industry



Moddable XS

Uses **Compartments** for safe end-user scripting of IoT products




MetaMask Snaps

Uses **LavaMoat** to sandbox plugins in their crypto web wallet



Agoric Zoe

Uses **Hardened JS** for writing smart contracts and Dapps



Google Caja

Uses **taming** for safe html embedding of third-party content



Mozilla Firefox

Uses **membranes** to isolate site origins from privileged JS code







Salesforce Lightning

Uses **realms** and **membranes** to isolate & observe UI components

Summary


This Lecture: Recap

- Part I: **why module isolation** is critical to modern JavaScript applications
- Part II: the **Principle of Least Authority**, by example
- Part III: safely composing modules using **least-authority patterns**



The take-away messages

- Modern applications are **composed from many modules**.
- You can't trust them all (**software supply chain attacks**)
- Apply the “principle of least authority” to **limit trust**.
 - Step 1: **Isolate modules** (Hardened JS & Lavamoat)
 - Step 2: Let modules **interact with “least authority”** (using reusable programming patterns)
- Understanding these patterns is **important in a world of > 2,000,000 NPM modules**.
- Even more critical in the emerging “**Web3**” where code can access valuable digital assets (think: tokens, NFTs, ...)



Designing “least-authority” JavaScript apps

Tom Van Cutsem
DistriNet KU Leuven

Questions?
tom.vancutsem@kuleuven.be



tvcutsem.github.io



be.linkedin.com/in/tomvc



github.com/tvcutsem



x.com/tvcutsem



@tvcutsem@techhub.social

Further Reading

- Mark Miller, Ka-Ping Yee, Jonathan Shapiro, “Capability Myths Demolished”: <https://srl.cs.jhu.edu/pubs/SRL2003-02.pdf>
- Compartments: <https://github.com/tc39/proposal-compartments> and <https://github.com/Agoric/ses-shim>
- ShadowRealms: <https://github.com/tc39/proposal-realms> and github.com/Agoric/realms-shim
- Hardened JS (SES): <https://github.com/tc39/proposal-ses> and <https://github.com/endojs/endo/tree/master/packages/ses>
- Subsetting ECMAScript: <https://github.com/Agoric/Jessie>
- Kris Kowal (Agoric): “Hardened JavaScript” <https://www.youtube.com/watch?v=RoodZSIL-DE>
- Making Javascript Safe and Secure: Talks by Mark S. Miller (Agoric), Peter Hoddie (Moddable), and Dan Finlay (MetaMask): <https://www.youtube.com/playlist?list=PLzDw4TTug5O25J5M3fwErKlmrjOrqGikj>
- Moddable: XS: Secure, Private JavaScript for Embedded IoT: <https://blog.moddable.com/blog/secureprivate/>
- Membranes in JavaScript: tvcutsem.github.io/js-membranes and tvcutsem.github.io/membranes
- Caja: <https://developers.google.com/caja> (Capability-secure subset of JavaScript)
- Chip Morningstar, “What are capabilities”: <http://habitatchronicles.com/2017/05/what-are-capabilities/> (broad historical perspective)
- Why KeyKOS is fascinating: <https://github.com/void4/notes/issues/41> (sketches the early history of capabilities as used in operating systems)
- Neil Madden, “Capability-Based Security and Macaroons” https://freecontent.manning.com/capability-based-security-and-macaroons/#id_ftn3 (capabilities in REST APIs)

Acknowledgements

- Mark S. Miller (for the inspiring and ground-breaking work on Object-capabilities, Robust Composition, E, Caja, JavaScript and Secure ECMAScript)
- Marc Stiegler’s “PictureBook of secure cooperation” (2004) is a great source of inspiration for patterns of robust composition
- Doug Crockford’s “JS: the Good Parts” and “How JS Works” books provide a highly opinionated take on how to write clean, good, robust JavaScript code
- Kate Sills and Kris Kowal at Agoric for helpful comments on earlier versions of these slides
- The Cap-talk and Friam community for inspiration on capability-security and capability-secure design patterns
- TC39 and the es-discuss community, for the interactions during the design of ECMAScript 2015, and in particular all the feedback on the Proxy API
- The SES secure coding guide: <https://github.com/endojs/endo/blob/master/packages/ses/docs/secure-coding-guide.md>